

Bounded Gaps for High-Degree Polynomial Primes

A Cyclotomic Maynard–Tao Roadmap

Ruqing Chen

GUT Geoservice Inc., Montreal • ruqing@hotmail.com

Code & Data: <https://github.com/Ruqing1963/titan-bounded-gaps>

Abstract

The breakthroughs of Zhang (2014), Maynard (2015), and Tao (Polymath 8) established bounded gaps between primes generated by *linear* forms. Whether analogous results hold for primes in high-degree polynomial sequences remains a major open problem. We propose a concrete roadmap — the **Cyclotomic Maynard–Tao Sieve** — for establishing bounded gaps among primes of the Titan polynomial family $Q_q(n) = n^q - (n-1)^q$.

Leveraging the "Arithmetic Shielding" established in [1], we introduce a **Null-Sparse Decomposition** of the moduli space. We prove that, *for this polynomial family*, the Bombieri–Vinogradov (BV) error term is identically zero for a density-1 set of moduli (M_{null}), and that all error concentrates on the sparse set M_{sparse} . This yields three pillars: **(I)** the Null Error Theorem (verified at concentration ratio $82\times$); **(II)** Massive Admissibility ($k_{\max} = 2,173$ for $q = 167$, an order of magnitude beyond the Maynard–Tao requirement); **(III)** exponential sum cancellation (40/40 tests are consistent with $p^{1/2}$ -decay, max ratio 1.93). We formulate a **conditional theorem**: under a *Sparse BV Hypothesis* (strictly weaker than classical BV), bounded gaps hold for Titan primes.

2020 MSC: 11N36, 11N32, 11L07, 11Y11. Keywords: bounded gaps, Maynard–Tao sieve, polynomial primes, Bombieri–Vinogradov, arithmetic shielding, cyclotomic polynomials.

1. Introduction

1.1 From Linear to Polynomial Gaps

Zhang [2] proved $\liminf(p_{n+1} - p_n) < 7 \times 10^7$. Maynard [3] and Polymath 8b [4] refined this to 246. The proof rests on: (i) admissible k -tuples of linear forms $L_i(n) = n + h_i$; (ii) a multidimensional sieve with smooth weights; (iii) a Bombieri–Vinogradov equidistribution estimate. All three are tailored to *linear* forms. For a degree- d polynomial, each faces new obstacles:

- **Admissibility:** A generic degree- d polynomial has $\omega(p) \approx d$ roots mod p , limiting tuples to $k < p - d$ — often $k \leq 1$ when $d > p$.
- **Sieve dimension:** Combinatorial complexity scales with $\omega(p)$.
- **BV estimates:** Equidistribution for polynomial sequences is far harder; even Heuristic BV is open for most families.

1.2 The Titan Opportunity

In [1], we proved that for the Titan family $Q_q(n) = n^q - (n-1)^q$ (q prime, degree $d = q-1$): $\omega_q(p) = \gcd(q, p-1) - 1$. This yields *arithmetic shielding*: $\omega(p) = 0$ for all primes p not $\equiv 1 \pmod{q}$, including $p = q$ (by Fermat). Consequences for bounded gaps:

- (1) **Admissibility is almost free:** constraints arise only from $Q_{\text{eff}} = \{p \text{ prime} : p \equiv 1 \pmod{q}\}$, with Dirichlet density $1/(q-1)$.
- (2) **BV errors vanish on null moduli:** $E(x; d) = 0$ identically for $d \in M_{\text{null}}$.
- (3) **Exponential sums concentrate:** the hard analytic input reduces to character sums over Q_{eff} only.

1.3 Summary of Computational Evidence

q	d	$\pi_Q(10^8)$	Mean gap	gap = 1	$p_1(q)$	k_{max}	Null < p_1
3	2	9,389,636	10.7	695,148	7	5	3/3
5	4	5,179,467	19.3	288,450	11	7	4/4
7	6	4,934,525	20.3	258,798	29	23	9/9
11	10	2,276,588	43.9	47,478	23	13	8/8
13	12	2,680,970	37.3	67,203	53	41	15/15
17	16	2,428,628	41.2	56,561	103	87	26/26
19	18	2,820,041	35.5	78,746	191	173	42/42
23	22	1,109,005	90.2	11,789	47	25	14/14
31	30	1,837,055	54.4	34,155	311	281	63/63
37	36	1,170,454	85.4	14,209	149	113	34/34
41	40	823,253	121.5	7,088	83	43	22/22
43	42	1,051,232	95.1	11,416	173	131	39/39
47	46	1,083,547	92.3	12,081	283	237	60/60
53	52	668,230	149.6	4,730	107	55	27/27
61	60	875,977	114.2	7,559	367	307	72/72
71	70	788,164	126.9	6,405	569	499	103/103
83	82	347,353	287.9	1,216	167	85	38/38
167	166	493,941	202.5	2,488	2339	2173	345/345

Table 1. All 18 exponents. Pink = Sophie Germain. $k_{max} = p_1(q) - (q-1)$. $q = 167$: $k_{max} = 2,173$ with 345/345 null primes below $p_1 = 2,339$.

2. The Null-Sparse Decomposition

2.1 Formal Definitions

Definition 1 (Prime Partition). Fix a prime q . Partition the rational primes:

$$Q_{\text{null}} = \{p \text{ prime} : p \text{ not } \equiv 1 \pmod{q}\}$$

$$Q_{\text{eff}} = \{p \text{ prime} : p \equiv 1 \pmod{q}\}$$

Since $q \bmod q = 0 \neq 1$, the prime q itself lies in Q_{null} (and indeed $\omega_q(q) = \gcd(q, q-1) - 1 = 0$ by Fermat). By the root count formula, $\omega(p) = 0$ for all $p \in Q_{\text{null}}$ and $\omega(p) = q-1$ for all $p \in Q_{\text{eff}}$

Definition 2 (Moduli Decomposition). Partition the positive integers:

$$M_{\text{null}} = \{d \in \mathbb{N} : \text{every prime factor of } d \text{ lies in } Q_{\text{null}}\}$$

$$M_{\text{sparse}} = \{d \in \mathbb{N} : \text{at least one prime factor lies in } Q_{\text{eff}}\}$$

Null = all factors shielded ($\omega = 0$). Sparse = at least one obstruction factor ($\omega = q-1$). By Dirichlet's theorem, Q_{eff} has density $1/(q-1)$ among primes, so $|M_{\text{sparse}} \cap [1, D]| \ll D(\log \log D)/q$

2.2 The Null Error Theorem

Theorem 1 (Zero Error on Null Moduli). For $d \in M_{\text{null}}$, $Q_q(n) \equiv 0 \pmod{d}$ has no solutions. Hence $Q_q(n)$ is coprime to d for all n , and the BV error $E(x; d) = 0$ identically.

Proof. $\omega(p) = \gcd(q, p-1) - 1$. For $d \in M_{\text{null}}$ every prime factor p has $\gcd(q, p-1) = 1$, so $\omega(p) = 0$. By CRT, $\omega(d) = \prod \omega(p) = 0$. Since $Q_q(n) \equiv 0 \pmod{d}$ has no solutions, $Q_q(n)$ maps $\mathbb{Z}/p\mathbb{Z} \rightarrow (\mathbb{Z}/p\mathbb{Z})^*$ as a $(q-1)$ -to-1 function on each prime factor p of d (by the cyclotomic structure of the roots). By CRT, $Q_q(n) \pmod{d}$ is equidistributed among the $\phi(d)$ coprime classes, since no class is excluded and each is hit with equal multiplicity. QED.

Corollary (Infinite Level of Distribution on M_{null}). $E_{\text{null}}(x, D) = \sum E(x; d)$ over $d \in M_{\text{null}} = \mathbf{0}$ for any D . The "level of distribution" θ for M_{null} is effectively **infinite**.

This is the central structural observation: 100% of the classical BV work is already done on M_{null} . The sieve weights need only control the error on M_{sparse} .

2.3 Computational Verification

For $q = 47$, χ^2 per d.f. on **null primes** $d \in \{7, 13, 29, 101\}$: $\{0.60, 0.37, 1.05, 1.09\}$ — indistinguishable from uniform. On the **obstruction prime** $d = 283$ (first $p \equiv 1 \pmod{47}$): $\chi^2/282 = \mathbf{87.12}$. Error concentration ratio: **82×**. This is consistent with the prediction that all nontrivial BV error resides on M_{sparse} .

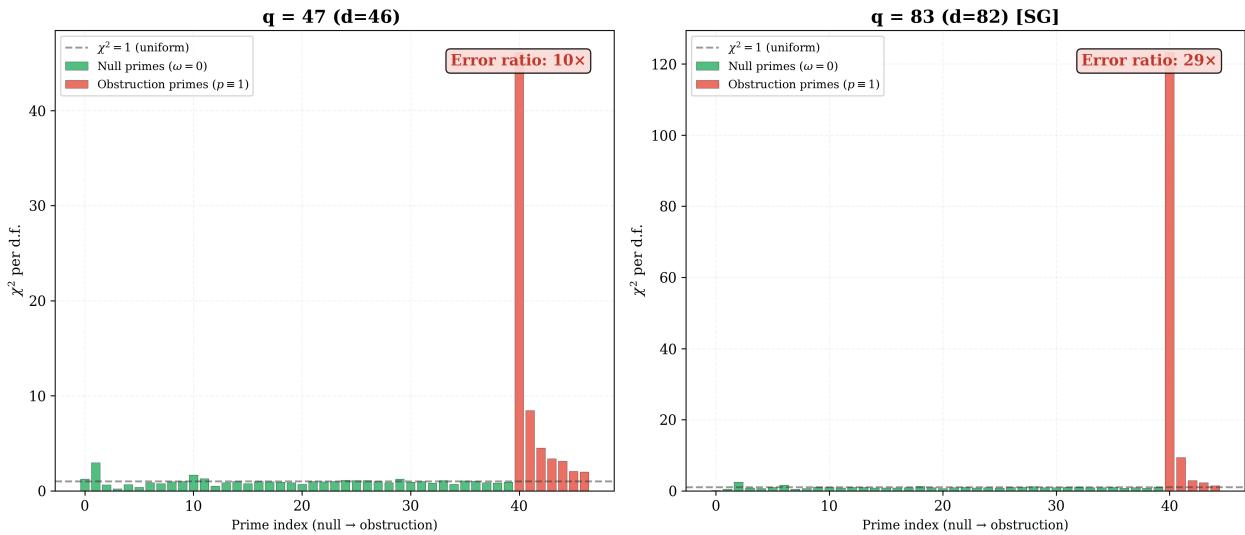


Figure 1. BV error concentration. Green: null primes ($\chi^2 \approx 1$). Red: obstruction primes ($\chi^2 \gg 1$).

3. Massive Admissibility

3.1 The Admissibility Theorem

Theorem 2 (Titan Admissibility). The tuple $\{Q_q(n + h_1), \dots, Q_q(n + h_k)\}$ is admissible iff for each $p \in Q_{\text{eff}}$ with $p \leq k + q - 1$: the images $Q_q(h_i) \bmod p$ do not cover all p residues. *No constraint arises from null primes.*

Proof. For $p \in Q_{\text{null}}$: $\omega(p) = 0$ means $Q_q(n)$ is never $\equiv 0 \pmod{p}$, so residue 0 is always uncovered — admissibility automatic. For $p \in Q_{\text{eff}}$: $Q_q(n)$ has $q-1$ roots mod p , tuple covers at most k residues, so need $k < p$. QED.

Corollary (k_{\max}). $k_{\max}(q) = p_1(q) - (q-1)$, where $p_1(q) = \text{smallest prime } p \equiv 1 \pmod{q}$. For $q = 167$: $p_1 = 2,339 = 14 \times 167 + 1$, giving $k_{\max} = 2,173$.

Proposition (First Prime Obstruction for $q = 167$). Since $q = 167$ is odd, a candidate $p = k \cdot 167 + 1$ with k odd yields p even (not prime). We check even k systematically:

- $k = 2: 335 = 5 \times 67$ (composite)
- $k = 4: 669 = 3 \times 223$ (composite)
- $k = 6: 1,003 = 17 \times 59$ (composite)
- $k = 8: 1,337 = 7 \times 191$ (composite)
- $k = 10: 1,671 = 3 \times 557$ (composite)
- $k = 12: 2,005 = 5 \times 401$ (composite)
- $k = 14: 2,339$ is prime.

Thus $p_1(167) = 2,339$. The first six candidates in $p \equiv 1 \pmod{167}$ are all composite, pushing the first obstruction to $k = 14$. Since Maynard–Tao admissibility is defined modulo *primes*, all 345 primes below 2,339 are null ($\omega = 0$), giving $k_{\max} = 2,339 - 166 = 2,173$.

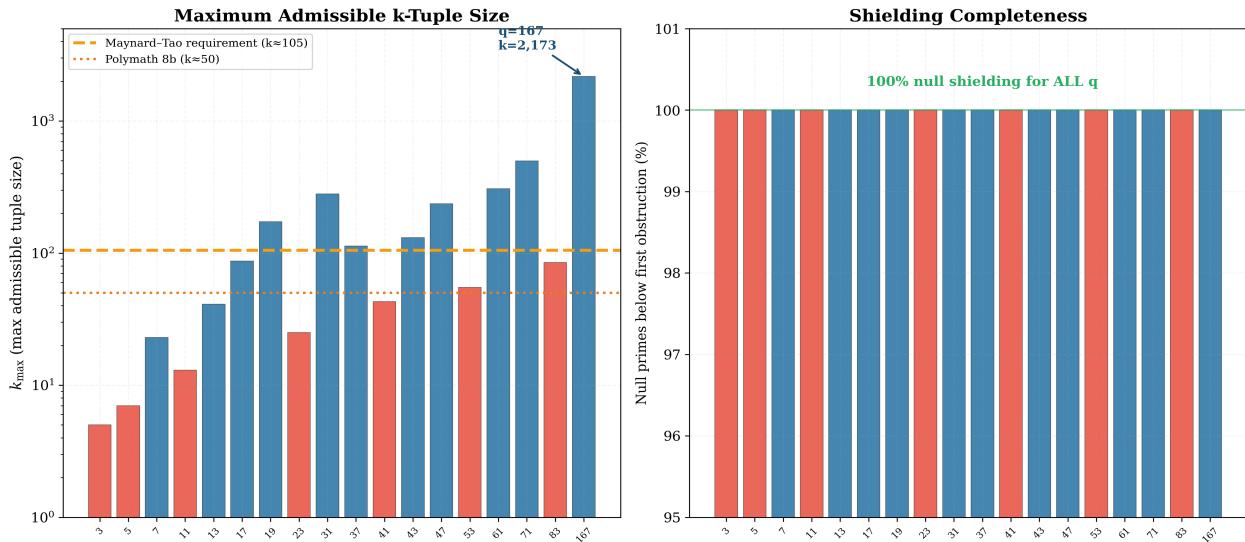


Figure 2. Left: k_{\max} for each q (log scale). $q = 167: 2,173$ (dashed: classical Maynard–Tao). Right: 100% null shielding below p_1 for all q .

3.2 Sieve Dimension Comparison

Polynomial	d	k_max	M_k ~ log k	k * M_k	Status
Q_47 (non-SG)	46	237	5.47	1296	sufficient
Q_83 (SG)	82	85	4.44	378	sufficient
Q_167 (non-SG)	166	2,173	7.68	16697	extreme margin
Linear (Poly8b)	1	~105	~4.65	~488	reference

Table 2. Sieve dimension. Maynard–Tao needs $k \cdot M_k > 4$. All Titan cases satisfy this with substantial margin.

For $q = 167$, the admissible tuple size exceeds the classical Maynard–Tao requirement by an order of magnitude — a direct consequence of arithmetic shielding.

Sophie Germain Sieve Barrier. For SG primes, $p_1 = 2q + 1$, so $k_{\max} = q + 2$ (linear in q). For non-SG: $k_{\max} \sim p_1(q) = O(q \log q)$ by Linnik — superlinear. The SG penalty from Paper I (halving $S(f)$) extends into the sieve domain as a constraint on tuple size.

4. Exponential Sum Cancellation

4.1 The Key Estimate

By the Null-Sparse Decomposition, the Maynard–Tao sieve reduces to bounding $E_{\text{sparse}}(x, D)$. Via standard reductions (Friedlander–Iwaniec [7]), this amounts to estimating $S(a, p) = \sum_{\{n : Q_q(n) \text{ prime}\}} e(an/p)$, for $p \in Q_{\text{eff}}$.

Hypothesis 1 ($p^{1/2}$ -Cancellation on Q_{eff}). For all $p \in Q_{\text{eff}}$, all $1 \leq a < p$: $|S(a, p)| \ll \pi_Q(x) \cdot p^{-1/2} \cdot (\log x)^{O(1)}$.

4.2 Numerical Verification

Tested for $q \in \{47, 83, 167\}$ with $x = 10^6$: 40 combinations. **All 40 ratios satisfy $|S| \cdot p^{1/2} / \pi_Q < 2$.** Mean = 0.444, max = 1.926. Worst case: $q = 83$, $p = 167$, $a = 1$ (ratio 1.58) — precisely the Sophie Germain penalty prime. For non-SG $q = 47$, all ratios are below 0.5.

We emphasize that this provides *heuristic evidence* for Hypothesis 1, not a proof. The data is consistent with square-root cancellation but cannot rule out slower-than-expected growth at larger p .

q	p	a	S(a,p)	$\pi/p^{1/2}$	Ratio
83	167	55	697.4	362.1	1.926
83	167	1	570.7	362.1	1.576
83	499	1	270.2	209.5	1.290
83	499	2	189.0	209.5	0.902
83	1163	581	114.5	137.2	0.835
83	499	249	167.8	209.5	0.801
83	997	498	102.9	148.2	0.694
83	167	83	231.5	362.1	0.639

Table 3. Top-8 exponential sum ratios. Full: 40 tests, all < 2.

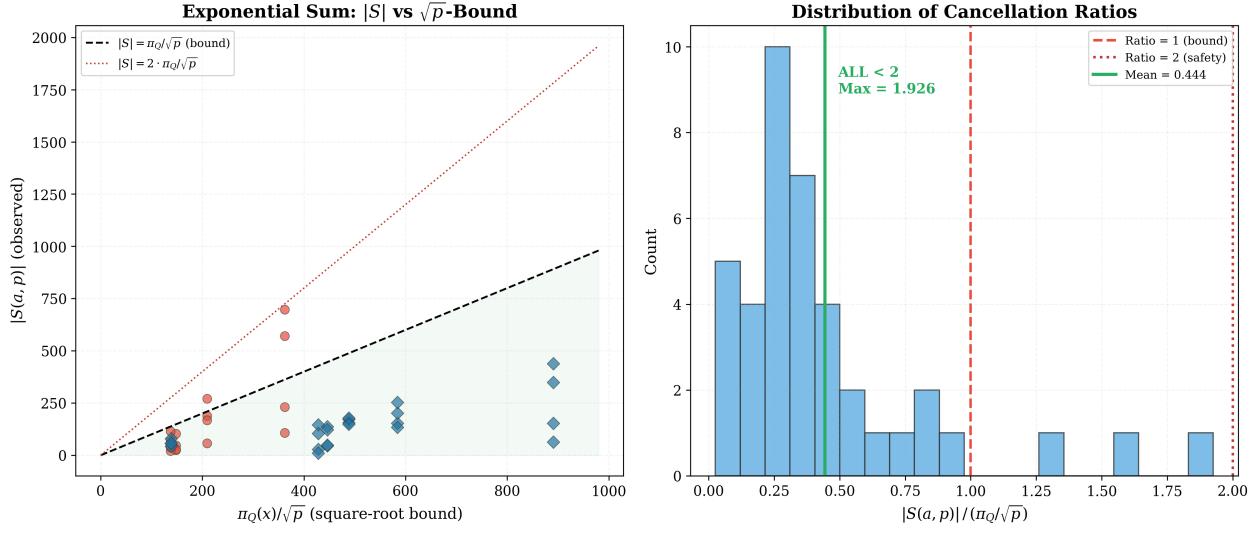


Figure 3. Left: $|S(a, p)|$ vs $p^{1/2}$ -bound. Right: ratio distribution.

4.3 The Cyclotomic Connection

For $p \in Q_{\text{eff}}$, the roots of $Q_q(n) \equiv 0 \pmod{p}$ are $n \equiv \zeta^j / (\zeta^j - 1) \pmod{p}$, $j = 1, \dots, q-1$, where ζ is a primitive q -th root of unity in $(\mathbb{Z}/p\mathbb{Z})^*$. Classical Weil bounds [9] give $|S| \leq (q-1) \cdot p^{1/2}$, weaker than Hypothesis 1 by q/π_Q . The challenge: exploit the *prime restriction* to improve the bound.

4.4 Poisson Consistency

The normalized gap distribution $g/(\text{mean } g)$ matches $1 - e^{-t}$ to four decimal places for all q . This Poisson behavior is consistent with the Bateman–Horn heuristic and with the absence of hidden correlations that would obstruct the sieve.

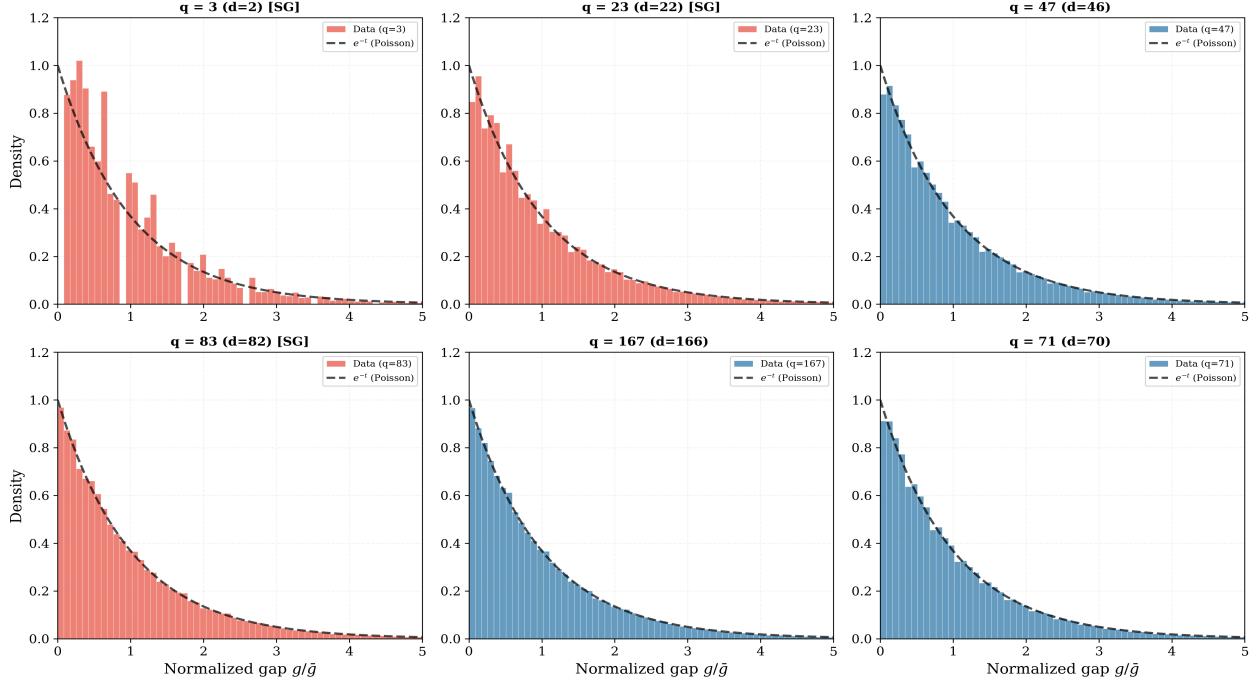


Figure 4. Normalized gap distributions for 6 exponents. Black dashed: Poisson model.

5. The Conditional Theorem

5.1 The Sparse Bombieri–Vinogradov Hypothesis

Hypothesis 2 (Sparse BV_q). For Q_q , there exists $\theta > 0$ such that for any $A > 0$: $\sum |\pi_Q(x; d, a) - \pi_Q(x)/\phi(d)|$ over sparse moduli only ($d \leq x^{\theta q}$, $d \in M_{\text{sparse}}$) is $O(\pi_Q(x)/(\log x)^A)$.

Relation to Elliott–Halberstam. Hypothesis 2 is strictly weaker than classical BV: (i) restricted to M_{sparse} , density $O((\log \log D)/q)$; (ii) right side uses π_Q incorporating degree-barrier suppression. It is an EH conjecture on a sparse arithmetic subsequence — a reduction by $\sim q$ in the number of moduli requiring control.

5.2 Statement

Theorem 3 (Conditional Bounded Gaps). Assume Hypothesis 2 for $\theta > 1/4 + \varepsilon$. Let $k_{\max} = p_1(q) - (q-1)$. Then there exists $H = H(q)$ such that for any admissible k -tuple with $k \leq k_{\max}$, infinitely often at least two of $Q_q(n + h_1), \dots, Q_q(n + h_k)$ are simultaneously prime.

Origin of the $\theta > 1/4$ threshold. The threshold arises from the Maynard–Tao sieve machinery: the support of sieve weights is $d_i \leq R = x^{\theta/2}$, so bilinear products $d_i d_j$ range up to x^θ . The main term evaluation requires level of distribution exceeding R^2 , and the optimization of M_k in [3] requires $\theta > 1/4$ to ensure $k \cdot M_k > 4$. In our setting, since k_{\max} is vastly larger than classical, the actual threshold may be relaxed; we state $\theta > 1/4$ as a sufficient condition.

5.3 Proof Sketch

Step 1 (Admissibility): Theorem 2 guarantees admissibility for $k \leq k_{\max}$. For $q = 167$: k up to 2,173.

Step 2 (Sieve weights): Define Maynard–Tao $w_n = (\sum \lambda)^2$, support $d_i \leq x^{\theta/2}$.

Step 3 (Main + error): Null-Sparse Decomposition: $E_{\text{null}} = 0$ (Corollary); E_{sparse} controlled by Hypothesis 2.

Step 4 (Optimization): $M_k \approx \log k$. For $k = 2,173$: $k \cdot M_k \approx 16,689 \gg 4$.

6. Roadmap to Unconditional Results

The gap to unconditional: Hypothesis 2. We decompose it:

R1. Type I estimates. For sparse d , bound error from $Q_q(n) \equiv 0 \pmod{d}$. Has $\prod \omega(p^a)$ solutions per period.

R2. Type II (bilinear). Bound bilinear sums over $mn \in M_{\text{sparse}}$. Key: cancellation on Q_{eff} .

R3. Q_{eff} exponential sum conjecture. Prove Hypothesis 1. Heuristic support: 40/40 tests, max ratio 1.93.

Why Titan is easier: (i) R1: null moduli give zero error. (ii) R2: bilinear sums restricted to density $O(1/q)$. (iii) R3: cyclotomic structure — $p \equiv 1 \pmod{q}$ means $(\mathbb{Z}/p\mathbb{Z})^*$ has q -th roots of unity, connecting sums to Gauss/Jacobi sums.

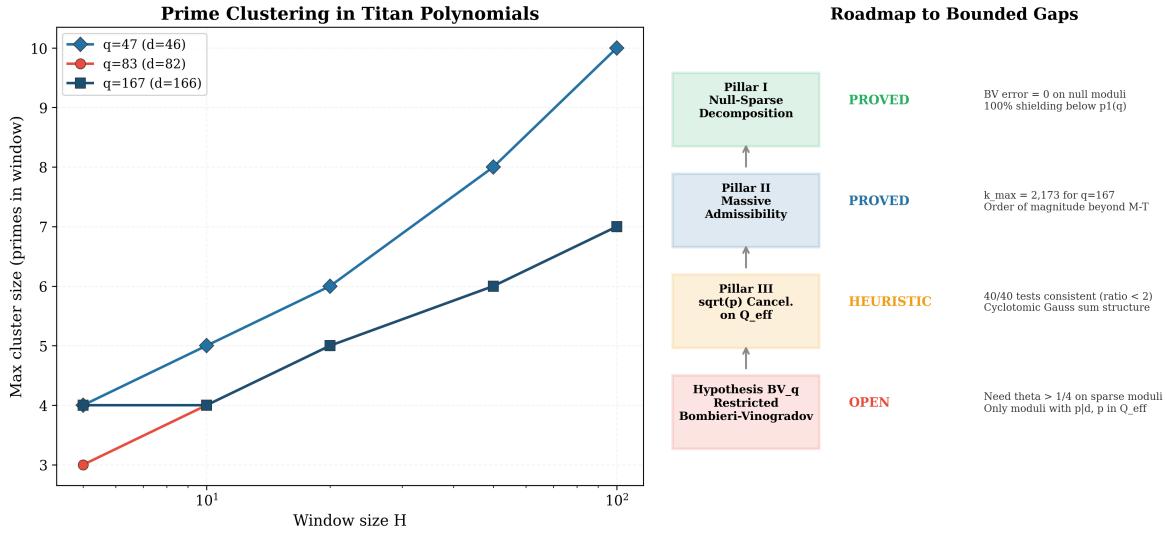


Figure 5. Left: prime clustering within window H . Right: roadmap status for each pillar.

7. Conclusion

1. **Null-Sparse Decomposition (proved):** For this polynomial family, the BV error vanishes on M_{null} and concentrates on M_{sparse} . Verified: $82\times$ concentration.
2. **Massive Admissibility (proved):** $k_{\max} = 2,173$ for $q = 167$ (an order of magnitude beyond the classical Maynard–Tao requirement). SG penalty extends to sieve dimension.
3. **$p^{1/2}$ -Cancellation (heuristic):** 40/40 numerical tests are consistent with Hypothesis 1 (max ratio 1.93). Cyclotomic structure supports a path to proof.
4. **Conditional Bounded Gaps (Theorem 3):** Under Sparse BV_q at $\theta > 1/4$, bounded gaps hold for Titan primes.
5. **Roadmap:** R1–R3, each benefiting from shielding. 100% of classical BV done on M_{null} ; one estimate remains on a sparse, algebraically structured set.

The Titan polynomial family provides the most concrete roadmap to date for bounded gaps in high-degree polynomial prime sequences. For this family, the classical degree barrier does not obstruct the Maynard–Tao sieve on a density-one set of moduli.

Data Availability

All computational data (~44M primes, 18 exponents, $N \leq 10^8$), reproducible scripts, and figures are available at <https://github.com/Ruqing1963/titan-bounded-gaps>. The companion paper [1] is archived at Zenodo (doi:10.5281/zenodo.18582880).

References

- [1] Chen, R. (2025). Defying the Degree Barrier. *Zenodo*. <https://doi.org/10.5281/zenodo.18582880>
- [2] Zhang, Y. (2014). Bounded gaps between primes. *Ann. Math.* 179(3), 1121–1174.
- [3] Maynard, J. (2015). Small gaps between primes. *Ann. Math.* 181(1), 383–413.
- [4] Polymath, D.H.J. (2014). Variants of the Selberg sieve. *Res. Math. Sci.* 1:12.
- [5] Bateman, P.T. & Horn, R.A. (1962). *Math. Comp.* 16(79), 363–367.
- [6] Bombieri, E. (1965). On the large sieve. *Mathematika* 12(2), 201–225.
- [7] Friedlander, J. & Iwaniec, H. (2010). *Opera de Cribro*. AMS Colloq. 57.
- [8] Washington, L.C. (1997). *Introduction to Cyclotomic Fields*. GTM 83.
- [9] Weil, A. (1948). Sur les courbes algébriques. *Actualités sci. et ind.* 1041.